

A hand is shown holding a glowing digital sphere. The sphere is composed of a network of nodes and lines, with a central core that resembles a circuit board. Various icons are visible on the sphere, including a location pin, a camera, a shield with a checkmark, a padlock, and a fingerprint. The background is dark blue with a faint network pattern.

Optimizing AI Models for Embedded Systems

Surendra Panpaliya



Agenda



AI Model Optimization Techniques



Techniques for Quantization,



Pruning, and Compression



Trade-offs between model size,



accuracy, and performance



in embedded systems



Agenda



Hardware Accelerators



Using GPUs, TPUs, and




other accelerators in embedded systems



Selecting appropriate hardware



for different AI tasks



Agenda



Power Management and Efficiency



Techniques for reducing power consumption in



AI applications on embedded devices



Balancing performance and



power efficiency for longer battery life

Agenda

Edge AI Frameworks and Tools

Overview of Edge Impulse,

TensorFlow Lite, Other frameworks for edge AI


Comparing different frameworks

for embedded AI development

Agenda



Hands-on Lab



Optimizing and
Deploying AI Models




on Embedded Systems

AI Model Optimization Techniques



Techniques for Quantization,



Pruning, and Compression



Trade-offs between model size,



Accuracy, and performance in
embedded systems


AI Model Optimization Techniques



REFER TO METHODS



USED TO IMPROVE
MODEL
PERFORMANCE,



EFFICIENCY, AND
ROBUSTNESS



DURING TRAINING
AND INFERENCE.


Model Optimization Techniques




ARE ESSENTIAL FOR
DEALING



WITH ISSUES SUCH AS
OVERFITTING,



UNDERFITTING, SLOW
CONVERGENCE, AND



INEFFICIENT USE OF
COMPUTATIONAL
RESOURCES.

Overfitting

Occurs when a model learns

the details and noise in the training data

to such an extent that

it negatively impacts

its performance on new data.

Overfitting

The model becomes too complex and

overly sensitive to the training data,

capturing both the signal (relevant patterns) and

the noise (irrelevant details).

Underfitting

Underfitting occurs

when a model is too simple

to capture the underlying

patterns in the data.

Underfitting

Can happen when

the model has too few parameters or

when it's not given enough training time.

Underfitting

As a result,

the model performs poorly

both on the training data and

the test data.

Slow Convergence

Training of a


machine
learning model

progresses
very slowly

Slow Convergence



Takes a long time for
the model



to reach an optimal
solution or



a satisfactory level of
accuracy.



Quantization, Pruning, and Compression

Quantization, Pruning, and Compression



Essential strategies



for optimizing AI
models,



particularly when
deploying them



in resource-constrained
environments


Quantization, Pruning, and Compression



Reduce the
model's size,



Computational
requirements



Energy
consumption,



while maintaining
accuracy.

The slide features a solid orange background. A large white circle is centered on the page. The word "Quantization" is written in a black, sans-serif font in the center of the white circle. On the left side of the white circle, there is a dashed blue arc. On the bottom right edge of the white circle, there is a solid purple circle.

Quantization

Quantization

Reduces the precision

of the numbers representing

model parameters (such as weights and activations)

from 32-bit floating-point numbers

to lower-bit representations

(like 16-bit, 8-bit, or even lower).

Quantization

Leads to

smaller model
sizes

faster inference
times.

Post-Training Quantization

Convert

the model

after it has been

fully trained.

Post-Training Quantization

Most common
approach

because it
doesn't require
changes

to the training
process.

Dynamic Range Quantization

Only weights are quantized,

typically from

32-bit floating-point

to 8-bit integers.

Full Integer Quantization

Both

weights and activations

are quantized

to 8-bit integers.

Float16 Quantization

Weights are stored in

float16 format (half precision),

which reduces the size

without impacting performance.

Quantization



Example: TensorFlow Lite post-training quantization



```
import tensorflow as tf
```



```
converter = tf.lite.TFLiteConverter.from_saved_model('model')
```




```
converter.optimizations = [tf.lite.Optimize.DEFAULT]
```



```
tflite_model = converter.convert()
```

Pruning



Pruning



Reducing its size



computation complexity



without impacting



its performance.

Pruning

Not all parameters

contribute equally

to the model's predictions,

so redundant or less impactful
ones

can be removed.

Magnitude-based Pruning

Remove
weights

that are
smaller than

a predefined
threshold.

Magnitude-based Pruning



```
import tensorflow_model_optimization as tfmot
```

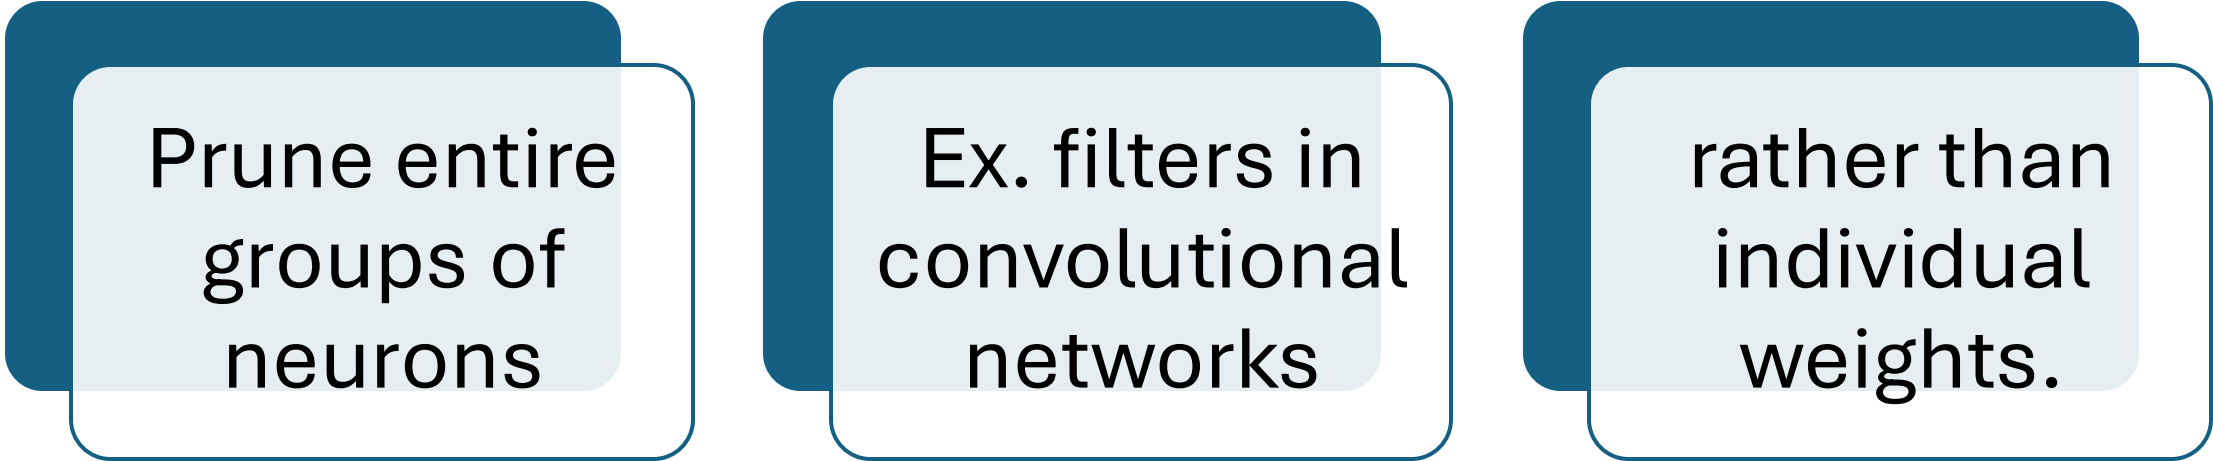


```
prune_low_magnitude =  
tfmot.sparsity.keras.prune_low_magnitude
```



```
model = prune_low_magnitude(model)
```

Structured Pruning



The diagram consists of three identical rectangular boxes arranged horizontally. Each box has a dark blue header and a light blue body. The text is centered within each box. The first box contains the text 'Prune entire groups of neurons', the second box contains 'Ex. filters in convolutional networks', and the third box contains 'rather than individual weights.'.

Prune entire
groups of
neurons

Ex. filters in
convolutional
networks

rather than
individual
weights.

Unstructured Pruning

Remove
individual
weights

in the model
based

on their
magnitude.

Compression



AIM TO REDUCE THE OVERALL
SIZE OF THE MODEL,



MAKING IT EASIER TO STORE
AND DEPLOY



WITHOUT LOSING ACCURACY.

Tools for Quantization, Pruning, and Compression



TensorFlow Lite



TensorFlow Model
Optimization Toolkit

OpenVINO



A toolkit for optimizing models,



including quantization and compression




for deployment on Intel hardware.

NVIDIA TensorRT



Supports model
optimization,



including quantization
and layer fusion,



specifically designed
for NVIDIA GPUs.

Trade-offs between model size, Accuracy, and performance



When deploying AI models in
embedded systems



there are critical trade-offs



to consider between



model size, accuracy, and
performance.

Model Size vs. Accuracy

Smaller models
tend to have

fewer
parameters and

lower
computational
complexity.

Model Size vs. Accuracy

Pruning and quantization

can be used to reduce model size


but may result in a slight decrease

in model accuracy

Trade-off Example



REDUCING A
MODEL'S SIZE BY 90%



(USING TECHNIQUES
LIKE PRUNING)



CAN DECREASE
ACCURACY BY 1-5%,



DEPENDING ON THE
TASK.

Trade-off Example



Quantized models (8-bit)



May sacrifice a small
percentage of Accuracy



compared to full
precision (32-bit) models

Consideration

For lightweight applications

like voice recognition on a smartphone

slight drops in accuracy

may be acceptable.

Model Size vs. Performance

Smaller models generally perform

better in terms of

latency (response time)

throughput (number of inferences
per second)

Model Size vs. Performance



Quantization can



reduce the model size



improve performance

Model Size vs. Performance

Pruned
models

have fewer
weights

to process
can

perform
better

in terms of
latency,


Trade-off Example



An 8-bit quantized
model



can run 3-4x faster than



its 32-bit floating-point
counterpart,



but might experience a
1-3% drop in accuracy.

Consideration

Latency-sensitive applications

real-time video processing or

autonomous driving

benefit more from

smaller, faster models

Optimizing AI Models

Hardware Accelerators

Optimizing AI Models Hardware Accelerators

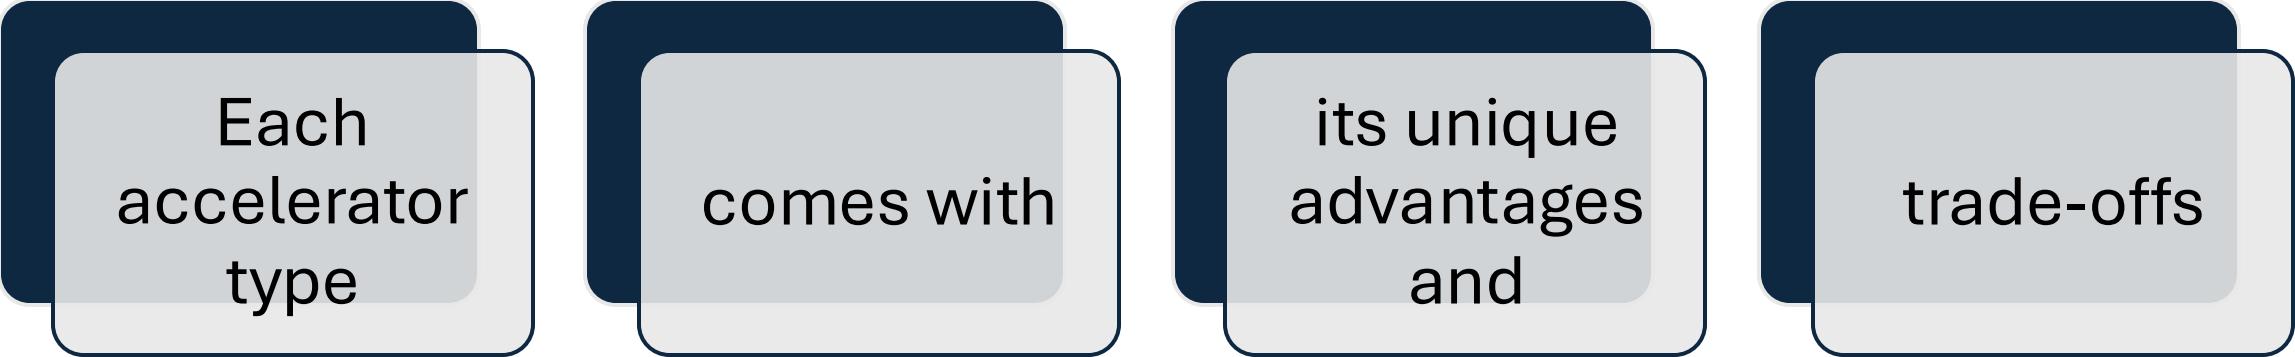
GPUs, TPUs, and specialized hardware

crucial for achieving high performance,

low power consumption, and

efficient model execution in embedded systems.

Optimizing AI Models Hardware Accelerators



Each
accelerator
type

comes with

its unique
advantages
and

trade-offs

Optimizing AI Models Hardware Accelerators




Selecting hardware



for different AI
tasks involves




considering factors
such as




computational
requirements,



Memory capacity,
power efficiency,
and cost.



General-Purpose vs. Specialized Hardware Accelerators



General-Purpose Hardware (CPUs)

Use Case:

General-purpose tasks,


lightweight AI models,

low-throughput applications.

Advantages



Highly versatile
and



can handle a wide
range of tasks,



including AI
inference and



general
computations.

Advantages



Suitable for
applications



where AI is just one
of many tasks



running on the
system.

Limitations

Lower parallelism compared

to specialized hardware (like GPUs or TPUs),


leading to slower inference speeds

for deep learning models.

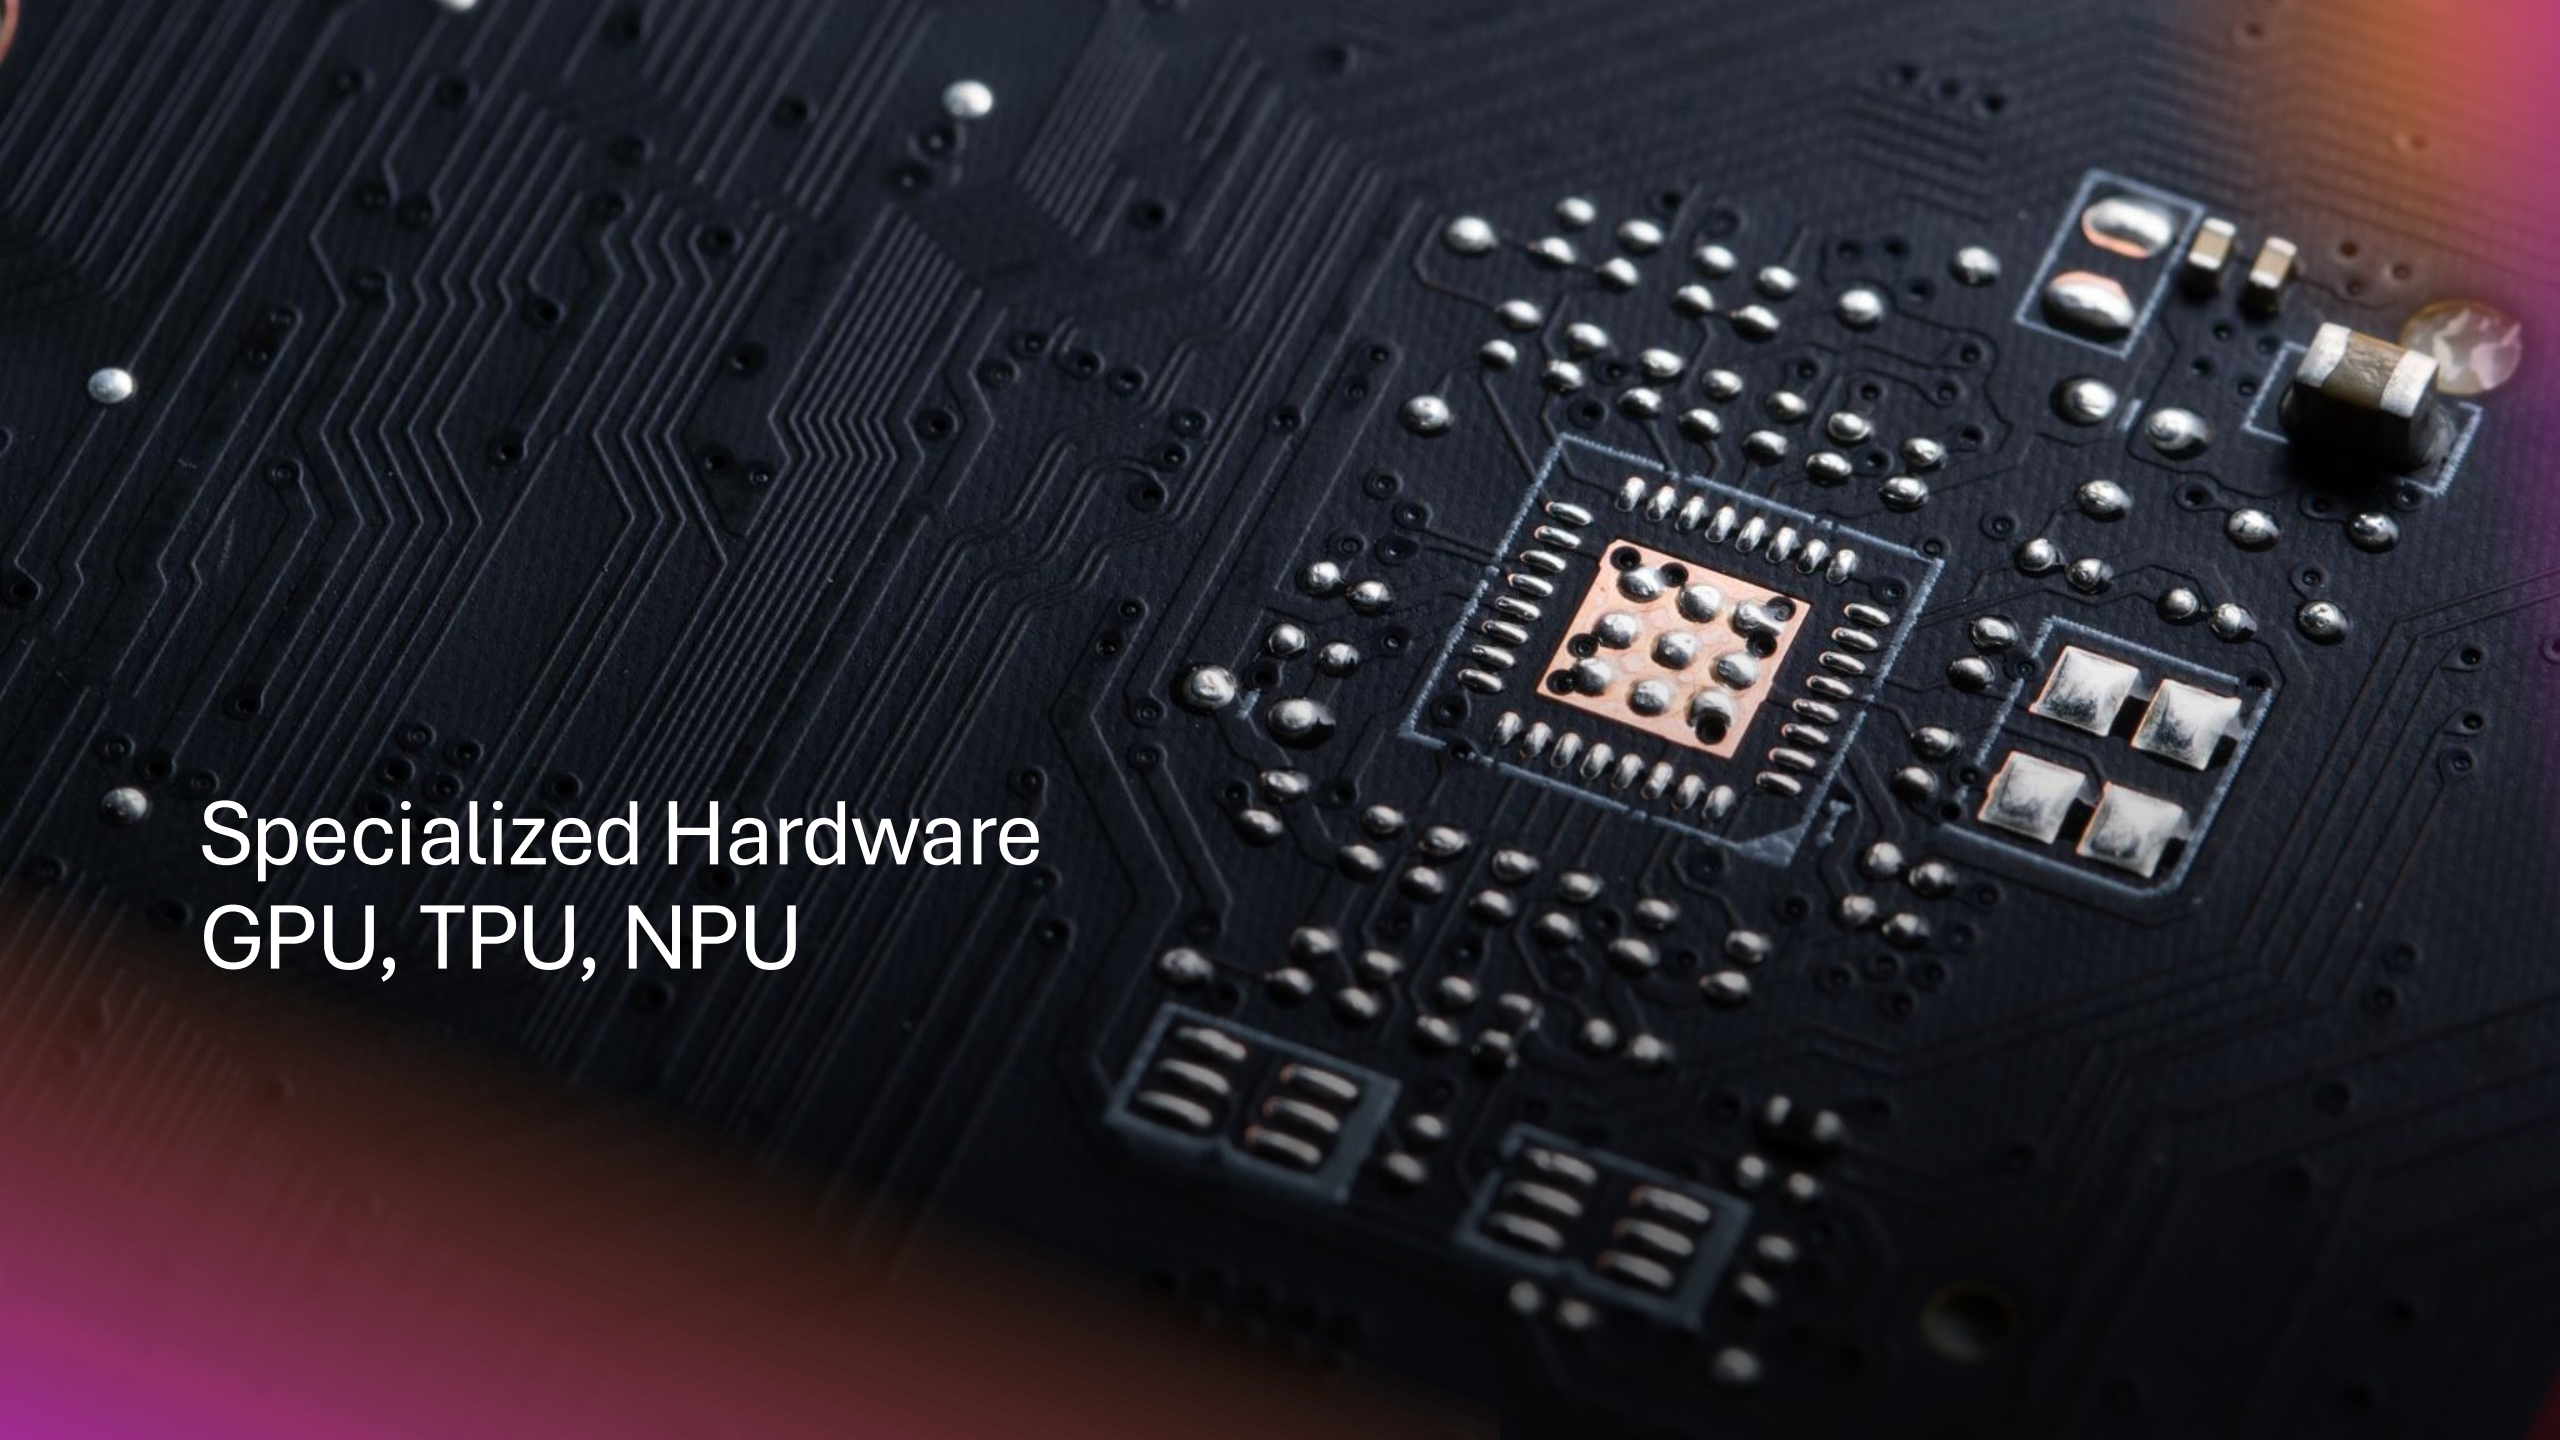
Limitations



Consumes more power for AI
tasks



compared to specialized
accelerators.



Specialized Hardware
GPU, TPU, NPU

Specialized Hardware



Optimized
for

High
throughput,

Parallel
computation

Efficient
execution of
AI models.

GPUs (Graphics Processing Units)

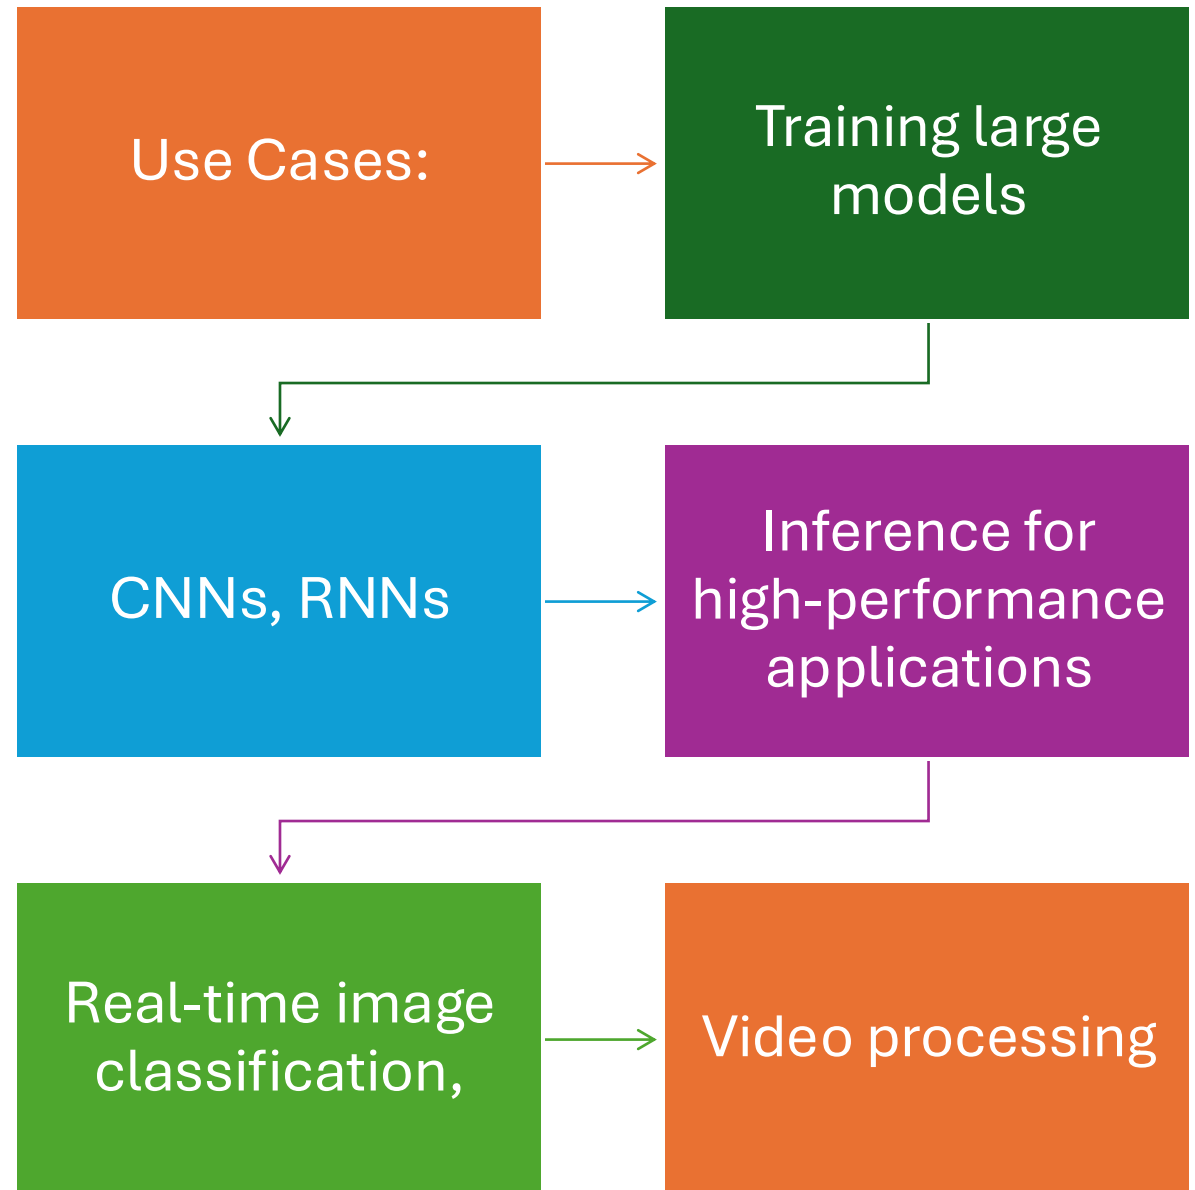
GPUs are designed for

parallel processing

well-suited for training and

inferencing deep neural networks.

GPUs (Graphics Processing Units)



Advantages




High parallelism:



Thousands of cores
handle



large-scale matrix and



tensor operations
efficiently.

Advantages

Flexibility:

Can be used for

both training and

inference tasks.

Advantages

Mature software ecosystem:

Popular deep learning frameworks like

TensorFlow, PyTorch

are highly optimized for GPUs.

Limitations




HIGH POWER
CONSUMPTION:



GPUS CONSUME
SIGNIFICANTLY MORE
POWER




LESS IDEAL FOR
LOW-POWER



EMBEDDED
SYSTEMS.

Best Fit



Embedded devices with



substantial power
budgets



Autonomous vehicles




Drones

Best Fit



Applications
requiring



real-time
processing of



high-dimensional
data



Video analytics,



Autonomous
driving

Examples of GPU Accelerators

NVIDIA Jetson:


A popular embedded AI platform for GPUs,

ideal for AI inference in robotics,


drones, and IoT devices.

(Jetson Nano, Jetson Xavier NX).


Examples of GPU Accelerators



Raspberry Pi 4 with
Google Coral USB
Accelerator:



Pairing a general-
purpose device like



Raspberry Pi with an
external accelerator



for GPU-level
performance.


TPUs (Tensor Processing Units)



TPUS ARE AI
ACCELERATORS



DESIGNED BY
GOOGLE,



OPTIMIZED FOR
TENSORFLOW.

TPUs (Tensor Processing Units)



Highly efficient at



running deep
learning models,




particularly for
inference.

Use Cases



Inference in large-scale applications:



TPUs excel at
deploying



AI models at scale

Use Cases

High-performance neural network tasks:

Models requiring tensor operations

CNNs, RNNs

work well with TPUs.

Advantages

Highly efficient for matrix/tensor operations:

TPUs are optimized for

matrix multiplications,

which are the backbone of

most neural networks.

Advantages



Low power consumption
compared to GPUs:



TPUs can deliver faster
inference



with lower power
requirements.

Advantages

Optimized for inference tasks:

TPUs are tailored more toward

high-speed inference

rather than training.

Limitations

Less flexible than GPUs:

TPUs are designed specifically for

tensor-based computations,

so they may not be as versatile

for non-tensor computations.

Limitations



Framework dependency:

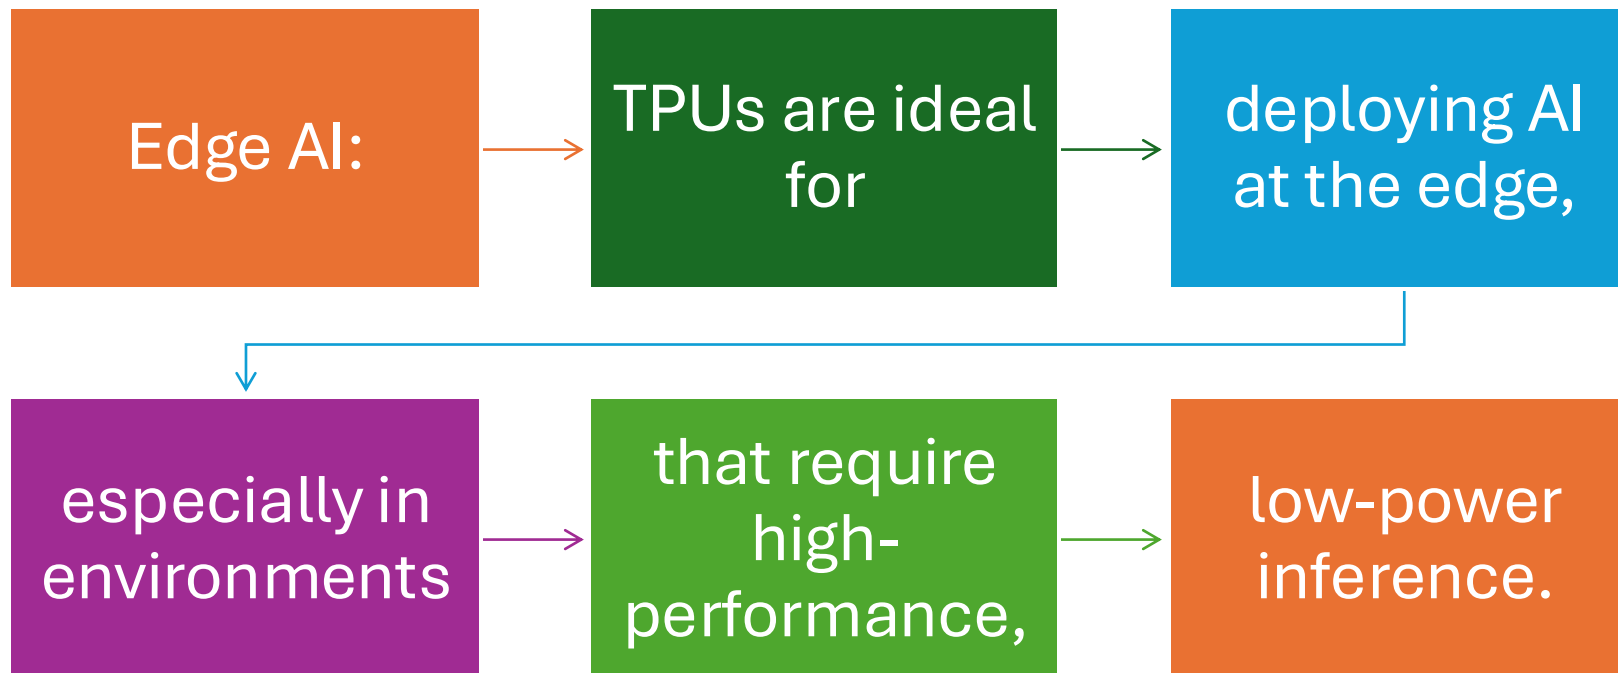


Best performance is achieved with TensorFlow,



limiting flexibility for other frameworks.

Best Fit:



Best Fit




Large-scale inference applications:



Suitable for tasks like real-time



Natural Language Processing (NLP) or



Computer Vision.

Examples of TPU Accelerators



Google Coral Edge TPU



An efficient AI
inference device



for edge computing
that delivers



high performance at
low power
consumption.

Examples of TPU Accelerators

Google Cloud
TPUs:

Used for large-
scale model
training and

inference tasks
in cloud
environments.

NPU (Neural Processing Units)



NPUS ARE SPECIALIZED
HARDWARE UNITS



DESIGNED FOR
ACCELERATING



NEURAL NETWORK
COMPUTATIONS.

NPU (Neural Processing Units)


NPU are highly optimized for

low-power embedded systems and

are commonly found in

smartphones and edge devices.


Use Cases



AI ACCELERATION IN
MOBILE DEVICES:



NPUS ARE
COMMONLY USED IN



SMARTPHONES,



TABLETS, AND SMART
CAMERAS

Use Cases

Real-time edge AI:

NPUs provide low-latency inference

for AI tasks at the edge,

without relying on cloud resources.

Advantages



Energy-efficient:



NPU's are designed
to consume



very little power
while delivering



high-performance
AI inference.

Advantages



Optimized for edge AI:



Perfect for low-power,



high-performance inference
tasks

Advantages




ON-DEVICE
INFERENCE:



CAN RUN MODELS
DIRECTLY ON THE
DEVICE



WITHOUT REQUIRING
CLOUD
CONNECTIVITY,



LEADING TO LOWER
LATENCY.

Limitations



Limited flexibility:



NPUUs are specialized and



may not handle



non-AI tasks efficiently.

Limitations

Smaller model capacity:

May struggle with large models

compared to more powerful

GPUs or TPUs.

Best Fit



Smartphones and mobile devices:



NPU's are typically integrated into



mobile SoCs for AI tasks like



face unlocking, camera enhancements, and



speech recognition.

Best Fit

Battery-powered IoT devices:

NPUs are ideal for low-power

AI inference in applications like


smart home devices or

wearable tech.

Examples of NPU Accelerators



Apple A-series and M-series
NPUs




Built into iPhones and iPads for
AI tasks.

Examples of NPU Accelerators



Huawei Kirin NPU:



Optimized for mobile AI
tasks



like real-time image
and video processing.



Surendra Panpaliya
Founder and CEO
GKTCS Innovations
<https://www.gktcs.com>

